Egg segmentation and tracking

Computer Vision Intern assignment, Nosh Robotics (Euphotic Labs). Tanmay Hutt, October 2026.

Code: https://github.com/tanmayhutt/egg-segmentation-tracking. Tracking video: https://drive.google.com/file/d/1DC3uYRb500AILF8IBrBE3OaTCb-4HBh4/view?usp=sharing.

1. Summary

Task	What was done	Result
1. SAM masks	SAM ViT-B automatic mode plus a measured egg filter; 14 of 110 photos needed one hand-placed click	110 masks, 138 eggs, all checked by eye
2a. Split	24 scene groups made by eye after three automatic grouping methods failed; whole groups dealt 4:1 with a seeded shuffle	Dev 88, eval 22, no scene on both sides
2b. Classical segmenter	Eight versions, tuned on dev only, about 15 side experiments, all logged	Dev mean IoU 0.371 (baseline) to 0.524 (final)
3. Evaluation	Final method run once on eval, scored against SAM	Eval mean IoU 0.472, precision 0.54, recall 0.64; baseline 0.188
4. Tracking	Classical segmenter per frame, Hungarian matching, occlusion memory, re-linking, scene-change reset	2 IDs for 2 eggs across the whole 32 s egg segment

Everything below was measured, not guessed. The complete trial-by-trial record, including the attempts that made things worse, is in `docs/experiment_log.md` in the repository.

2. The data

110 photos, all 640x640, exported from Roboflow as a classification dataset: one label per photo (`not_damaged`) and no masks, so Task 1 has to create the reference masks. Roboflow stretched every photo to a square, so egg shapes are slightly distorted. Camera metadata was stripped at export.

The photos are hard on purpose: eggs on a white towel, in ash, in steel bowls of water (three while boiling), under leaves, in hands, on a keyboard, in a pink bowl, on foil. Many scenes were shot several times from different angles, and the photo numbers do not follow scenes (the keyboard is #17, #18, #19 and #43). Roboflow's own train/valid/test split puts the same scene on both sides, so it was not used.



All 110 photos, sorted by number, with Roboflow's split tag

3. Task 1: reference masks with SAM

How SAM works. An image encoder turns the photo into an embedding (about 1 s on the M2 GPU). A prompt encoder takes a click, box or rough mask that says which object is meant. A mask decoder returns the mask, with up to three guesses and a confidence (predicted IoU) each. SAM knows what object boundaries look like but not object names, so it cannot be asked for "egg".

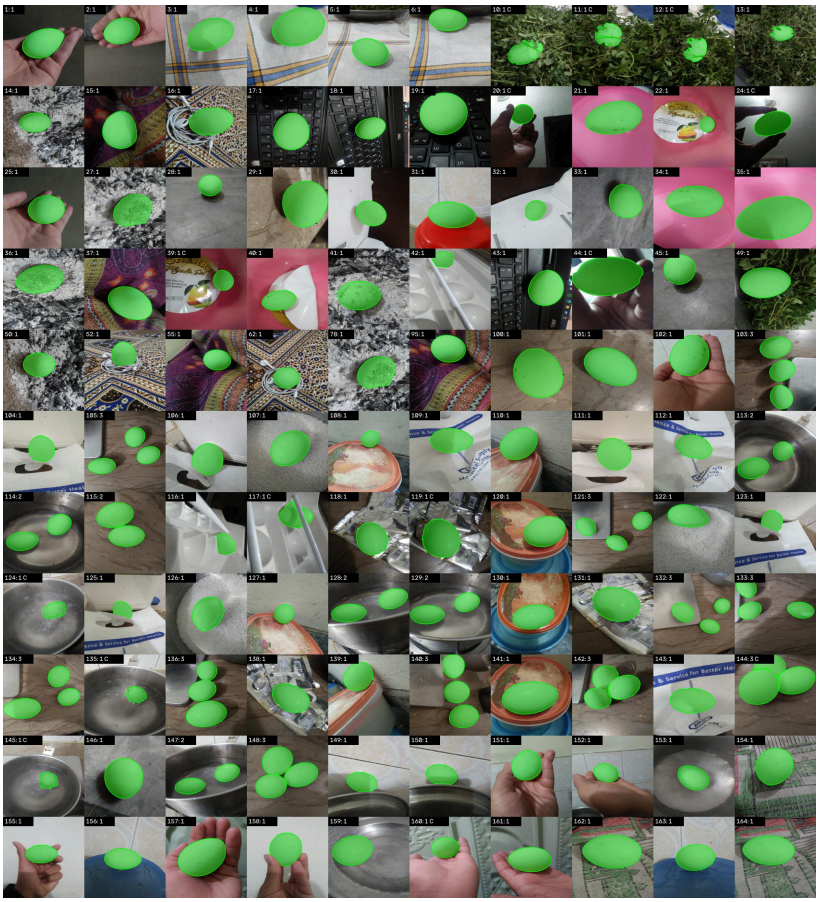
Automatic mode. With no clicks available, SAM's automatic mode places a 32x32 grid of points and returns a mask for everything: eggs, hands, keys, shadows, bowls (7 to 93 masks per photo). The slow step (`sam_candidates.py`, about 18 s per photo) saves every candidate once; the fast step (`sam_select.py`) picks the eggs in milliseconds, so the filter could be changed without re-running SAM.

Choosing which masks are eggs. Rules measured on 8 sample photos:

Rule	Why
0.8% to 60% of the photo	Drops specks and whole-background masks
Brightness (HSV V) 120 or more	Eggs measured 137+; shadows, keys and ash were darker
Saturation 45 or less	Eggs are nearly colourless
Solidity 0.95 or more (area / convex hull area)	An egg has no dents; a hand does
Ellipse fill 0.96 or more (area / fitted ellipse area)	Eggs 0.97+, empty egg-tray cups 0.90 to 0.95
One connected piece (95%+)	Rejects masks split by a rod or leaf
At least 25% of the biggest egg	Drops small egg-like things such as earbud heads
Not mostly inside a kept egg	SAM returns some eggs twice, and parts such as glare spots

Result. 96 photos fully automatic. 14 needed a hand-placed click, each listed with its reason in `config/sam_prompts.json`: leaves across the egg (solidity), dark backlit photos (brightness), eggs cut by the

frame (ellipse fill), a rod across the egg, and the three boiling-water photos where no clean mask existed. One click is ambiguous (glare spot, egg, or the bowl), so SAM returns three guesses and the most confident is kept; for the boiling-water photos the most confident guess was the whole bowl, so the smallest guess (the egg) is kept for those three only. 138 eggs in total: 93 photos with one egg, 6 with two, 11 with three. Every mask was checked by eye on the review sheet.



Task 1 result: every photo with its SAM egg mask. "C" marks photos fixed with clicks

Limits. ViT-B was used (ViT-H is 2.4 GB and slower; for a simple object ViT-B gave clean edges and no comparison was run). These masks are SAM's opinion plus a filter, not hand-drawn ground truth, so Task 3 measures agreement with SAM. The boiling-water masks are the least certain.

4. Task 2a: a split without leakage

Rules tuned on photo A and scored on photo B tell nothing if A and B are the same keyboard from two angles. So photos of one scene must land on one side. The standard way is to split by capture session, but the export carries no metadata, and the numbering does not follow scenes (20 of the 24 scenes are spread over non-adjacent numbers).

Three automatic groupings were tried and all failed. Each was tested on the known same-scene sets (keyboard #17/18/19/43, steel bowl #113/114/128/129): they must be closer to each other than to unrelated photos.

Fingerprint	Known same-scene pairs	Unrelated nearest neighbours	Why it fails
Perceptual hash (64 bits), Hamming distance	18 to 34 bits	16 to 24 bits	It fingerprints the layout of light and dark, which changes when the camera moves
HSV colour histogram, cosine distance	0.09 to 0.34	0.0 to 0.3	Nearly every photo is "white egg on grey or white", so histograms all look alike
SAM image embedding, averaged, cosine distance	0.003 to 0.018	0.00 to 0.02	Same, and averaging discards where things are

What was done. All 110 photos were grouped by eye into 24 scenes (`config/scene_groups.json`); when unsure, photos were merged, since over-merging is safe and under-merging leaks. `make_split.py` shuffles the groups with a fixed seed and deals whole groups to eval until it holds one fifth of the photos. The seed (21, chosen from 0 to 199) keeps eval's share of multi-egg photos (0.23 vs 0.15 overall) and of click-fixed photos (0.14 vs 0.13) close to the whole set: stratified sampling at scene level. Result: dev 88, eval 22, exactly 4:1. Eval scenes: striped towel, ash pile, steel bowl with water (including the three boiling photos), brown marble, red lid. Eval holds no hand photos and no beige-marble-counter photos; the split is frozen in `config/split.json`.



The 22 eval photos

5. Task 2b: the classical segmenter

"Classical" means no learning: hand-written rules turn pixel values into a mask. Every number in them was chosen by measuring on the 88 dev photos. Score: mean IoU against SAM's masks over dev.

Measurement before any rule. Per-photo median of egg pixels vs background pixels (10th to 90th percentile over dev): gray 168 to 209 vs 53 to 173, HSV V 172 to 215 vs 61 to 186, HSV S 4 to 24 vs 7 to 123, hue a coin flip. Brightness is the main signal and low saturation the second, but the ranges overlap (the brightest backgrounds reach the darkest eggs), so no single fixed threshold can work.

The final pipeline (v8), step by step.

1. Convert to HSV. Seeds are the brightest 15% of pixels that are also colourless (S <= 80).
2. Canny edges, thickened to 5 px, are removed from the seeds, so an egg and the surface it sits on start as separate pieces.
3. Distance transform: inside each seed blob keep only pixels at 70% or more of the blob's deepest point. A fat egg keeps a solid core; a thin bridge leaking into the surface is never deep.
4. Background seeds: pixels that are dark (darkest 30%) or strongly coloured (S > 60) are not egg.
5. Watershed floods from every seed until regions meet at edges. Each egg seed becomes a region.

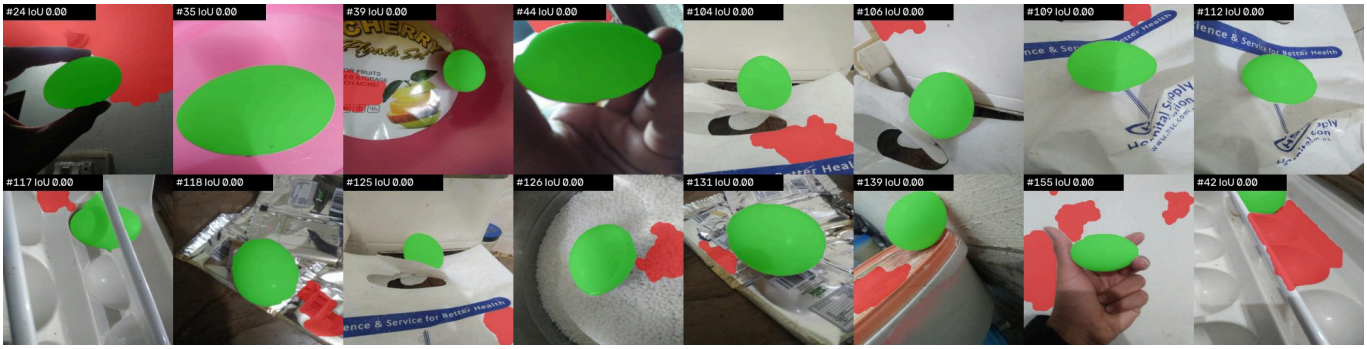
6. Each region is opened with a 31 px ellipse (erode, then dilate back) to cut necks into the surface, then smoothed to its outer outline.
7. Keep regions covering 2% to 60% of the photo with solidity >= 0.85 and ellipse fill >= 0.9. Drop regions under a quarter of the biggest kept one, and regions mostly inside a kept one.
8. Fallback: every photo in this dataset holds an egg, so if nothing passes, keep the single most egg-shaped region rather than return an empty mask.

Every version, with what happened.

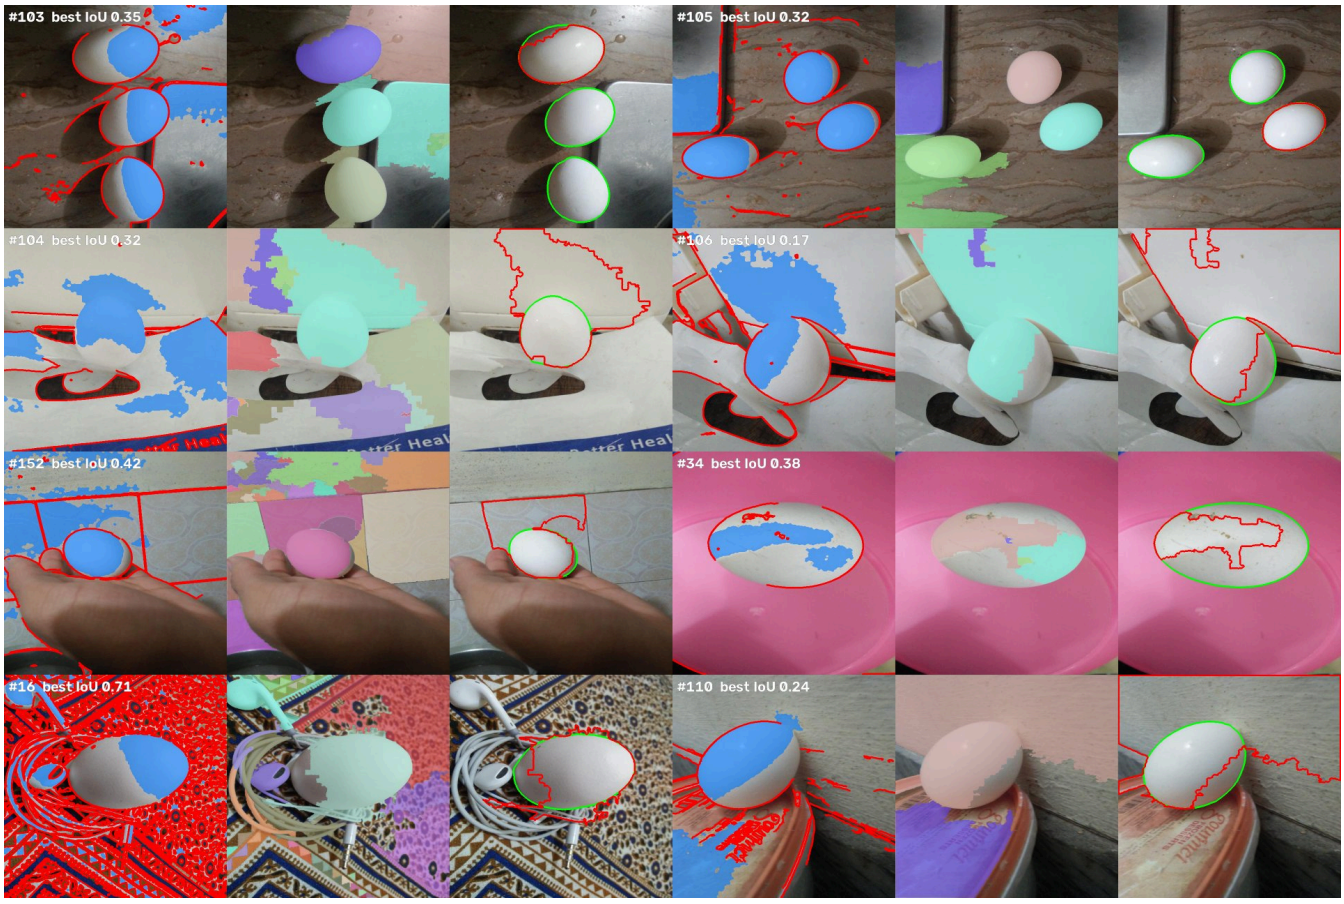
#	Change	Dev IoU	Under 0.5	What the result taught
v1	Grayscale, Otsu threshold, largest blob	0.371	61/88	Skin and white surfaces are bigger than the egg, so "largest" is wrong
v2	Add colourless (S <= 60), opening and closing 7 px	0.425	55/88	Hands in normal light fixed; steel, paper, white bags still merge with the egg
v3	Choose blobs by shape, not size	0.283	59/88	Worse, median 0. In 51 photos the egg was already merged into a bigger blob after thresholding; no later rule can split it. Lesson: measure each stage separately
side	Nine thresholding methods scored by "does the egg survive as its own blob"			Best plain threshold 45/88; watershed with edge-cut seeds 55/88. A 101 px local-contrast window scored 0 because it was smaller than the egg
v4	Watershed regions, then the v3 shape filter	0.350	55/88	Regions had a 0.67 per-egg ceiling, yet the pipeline scored 0.35: the selection step was the bottleneck
v5	Smooth regions; minimum area 2%; relative-size rule	0.465	46/88	Good egg regions have median area 9%, bad scraps 1%. Size separated them better than shape
v6	Distance-transform seeds	0.450	45/88	Eggs with a matching region 76 -> 86 of 109; mean dipped because the selection rules were tuned for the old regions. A 96-point grid over thresholds gave at most 0.477
v7	Coloured pixels as background seeds	0.451	42/88	Median 0.45 -> 0.62, but 42 photos still at exactly 0
v8	31 px opening to trim leaks; fallback	**0.524**	36/88	Zero-score photos 37 -> 15. Final

One measurement error is also in the log: the first "does the egg survive" test compared single regions to multi-egg masks, which capped the score at about 0.33 for three-egg photos. It was fixed to score egg by egg.

Where it works and fails on dev, by scene. Marble counter 0.83 (13 photos, including every touching-egg case), patterned cloths 0.83 to 0.84, keyboard 0.75, leaves 0.76. Hands in normal light 0.44 to 0.68. Pink bowl with a white sticker 0.36, foil 0.30, white bag 0.14, egg tray 0.10, backlit hands 0.00 (the egg is darker than the background). The common thread: a bright surface touching the egg with no visible edge between them. Brightness, colourlessness and edges are the only cues a classical method has, and all three are absent there.



The 16 worst dev photos for v8. Red: ours only. Green: SAM only. Yellow: both agree



Debugging sheet behind v6: seeds and edges (left), watershed regions (middle), best region vs SAM (right)

6. Task 3: evaluation on eval

Eval was run exactly once, after all tuning stopped (`scripts/evaluate.py`).

Version	Mean IoU	Median	Precision	Recall	Under 0.5	At 0
v1 baseline	0.188	0.168	0.206	0.605	20/22	7
v8 final	**0.472**	0.452	0.544	0.635	12/22	5
v8 on dev, for reference	0.524	0.611	0.605	0.681	36/88	15

Why these metrics. IoU (overlap over union) is the main score because it punishes both missing egg and grabbing background. Precision says what share of our mask is really egg; recall says what share of the egg we found. Pixel accuracy was not used: an egg covers about 5% of a photo, so an empty mask would score 95%.

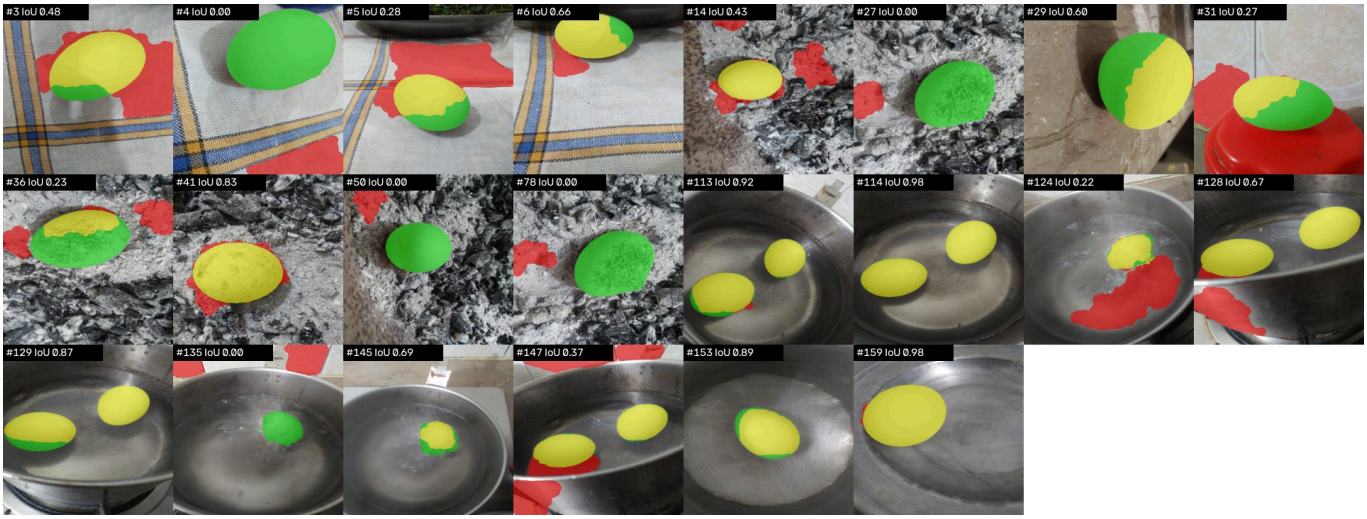
Reading the numbers. Dev to eval dropped by 0.05, the usual sign of mild over-tuning rather than a collapse. The baseline scored 0.19 on the same photos, so the tuning gained about 0.28 IoU on photos it never saw. Precision 0.54 means about half of what we call egg is not egg: extra bright regions (towel

patches, steel rims, ash). Recall 0.64 means a third of the egg is missed, mainly its shaded side, which falls below the brightness cut.

Per scene and failure analysis.

- Steel bowl (10 photos): 0.66. Clear water 0.87 to 0.98. Boiling water 0.00 and 0.22: steam and bubbles blur the outline, and SAM's own masks there were the least certain. Bright steel rims and reflections are kept as extra regions.
- Ash pile (6): 0.25, three at 0.00. The egg is dusted with ash and has the same grey as the pile: no brightness, colour or edge difference. The fallback keeps a light ash patch instead. A clean egg on ash scores 0.83.
- Striped towel (4): 0.36. White cloth passes every rule an egg passes; towel patches are kept beside the egg, and one frame-filling egg is missed entirely.
- Brown marble (1): 0.60. Red lid (1): 0.27, the bright tiled floor joins the egg region.

The eval set has no hand or marble-counter photos and does hold the two hardest scenes, so 0.47 is probably a slightly pessimistic estimate. It is reported as measured.



All 22 eval photos. Red: ours only. Green: SAM only. Yellow: both agree

7. Task 4: tracking in the video

The video ("Raw and Boiled Egg Spinning", 60 s, 24 fps, 1280x720) shows two white eggs on a dark table. A hand enters, spins each egg and carries one across the table. Title cards fill 0 to 10 s and 42 to 60 s. Output: `outputs/video/tracked.mp4` (Drive link on page 1) and `tracks.csv` (frame, time, ID, centre, area).

Adapting the segmenter to the frames (allowed by the brief), measured on 16 sample frames at 640x360, fallback off so title cards return nothing:

Change	Why	Effect
Minimum area 2% -> 0.4%	The two eggs measured 0.0202 and 0.0199 of the frame, one each side of the 2% cut	Both eggs found, but walls, shirt and table reflections also kept
Ellipse fill capped at 1.1	A rounded rectangle over-fills its fitted ellipse; eggs sit at 0.96 to 1.03	Small gain
Keep regions at least 85% as bright as the brightest kept one	Eggs measured 232 to 238; wall 141 to 150; reflections 137	Clean: 2 eggs when visible, 1 when the hand covers one, 0 on title cards

Photo results are unchanged.

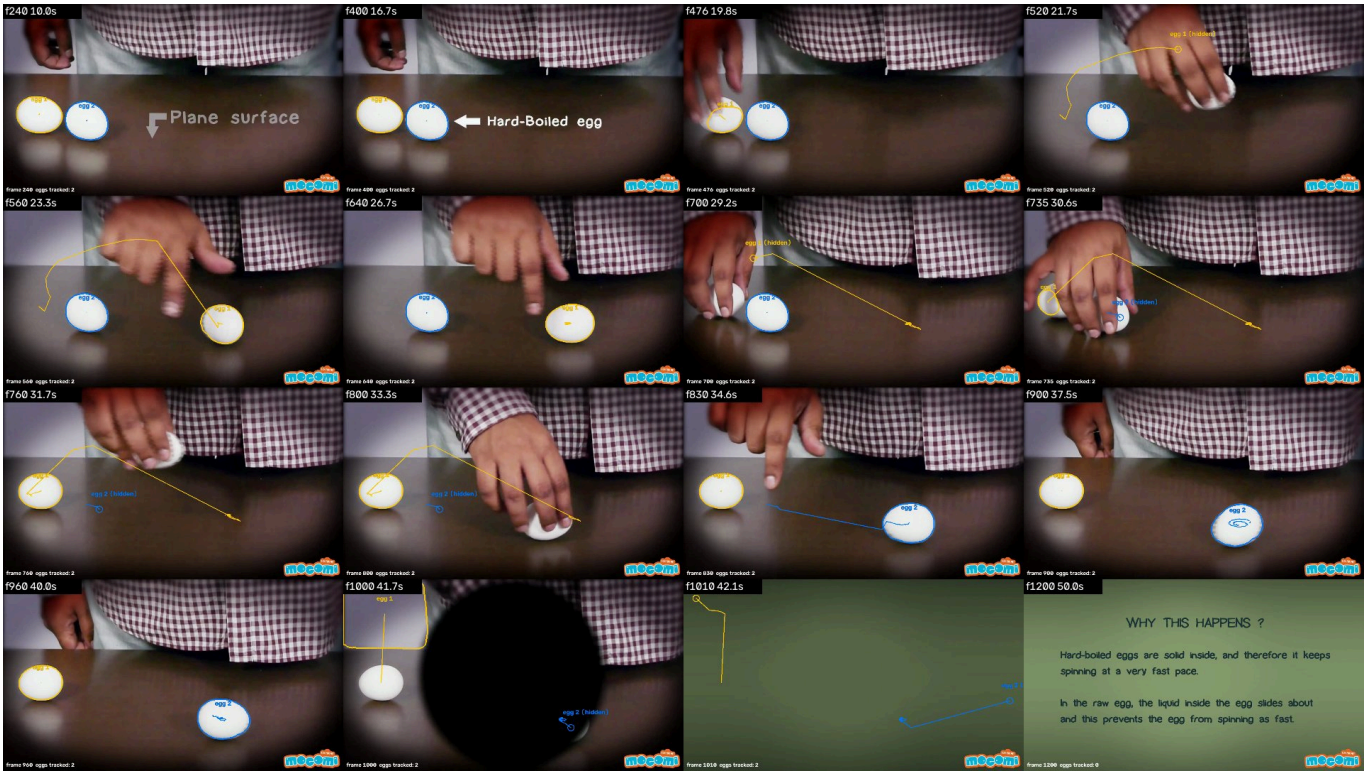
The tracker (`scripts/track.py`), per frame:

1. Detect eggs with the adapted v8. Each region gives a centre, an area and a contour.
2. Match detections to existing tracks by centre distance with the Hungarian algorithm (`scipy.optimize.linear_sum_assignment`), which solves the whole pairing at once so two eggs never both grab one detection. Matches over 60 px are refused.
3. A track with no match keeps its last position for up to 5 s (the hand hides a spinning egg for up to 3.7 s) and is drawn as a small circle marked "hidden".
4. A detection that matches nothing nearby is offered to a hidden track with no distance limit, if it is between half and double that egg's area. Eggs do not appear from nowhere: a new egg while one is missing is that egg, carried by the hand.
5. A new track needs 12 consecutive frames (0.5 s) before it gets an ID.
6. If the frame differs a lot from the previous frame or from the frame 0.5 s earlier, all tracks are dropped (hard cuts and fades).

Iterations.

#	Settings	IDs for 2 eggs	What went wrong
1	Memory 0.5 s, 3 frames to confirm	7	The hand carried an egg far while hiding it; it got a new ID on reappearing. A wrist and a fade each got an ID
2	Memory 3 s, 12 frames to confirm, re-link rule, hard-cut reset	3	One egg hidden 3.7 s by the spinning hand
3	Memory 5 s	2	During the iris fade the brightening wall was re-linked as egg 1, and tracks lingered onto the title card
4	Re-link needs similar area; fade detection	**2**	Egg 1 and egg 2 from 10.1 s to 41.6 s, 1281 detections, nothing on title cards

Limit. The two eggs look identical, so if both were hidden at once and swapped, the IDs would swap. That did not happen in this clip. SAM per frame was not an option: about 18 s per frame on this machine, 7 hours for the clip, against 40 s for the whole clip with the classical method.



Frames from the tracked video. Yellow: egg 1. Blue: egg 2. Lines are trails; a circle means hidden

8. What I would do next

- Recover the shaded side of each found egg with a second, lower brightness threshold applied only inside the fitted ellipse. This targets the recall loss and would not help the zero-score photos.
- For white-on-white and ash-on-ash scenes no pixel rule works; that is where a learned model earns its place. SAM found clean masks on 96 of 110 photos with no tuning at all.
- Group scenes with a scene embedding (CLIP or DINOv2) instead of by eye, so the split scales.
- Evaluate against hand-drawn masks on a subset, to check how far SAM itself is from the truth.

9. AI tools

An AI assistant (Claude) was used as a coding aid: drafting script code, helping debug, and explaining methods on request. The approach, the decisions at each step, the thresholds (all chosen from numbers measured on the dev photos), the checking of every result by eye, and the experiment log are my own work. No egg mask was produced by a language model.

Appendix: running the code

See `README.md`. In short: create a Python 3.12 environment from `requirements.txt`, place the dataset at `data/egg/`, the SAM weights at `weights/`, the video at `data/video/eggs.mp4`, then run `scripts/sam_candidates.py`, `sam_select.py`, `make_split.py`, `classical.py v8 dev`, `evaluate.py v8 eval`, `track.py` in that order. `build_report.py` rebuilds this PDF.